

<b>Course:</b>	CSCI 2050U: Computer Architecture I
<b>Topic:</b>	Debugging with gdb

## Overview

The purpose of this document is to give you skills in `gdb`, which is the most popular debugger on the Linux platform (as well as several other platforms). `gdb` is feature-packed, but its interface is entirely text-based.

## Part 1 – Getting Started

Let's get our assembly language program ready for debugging. You need to assemble project with the `-g` (and/or the `-ggdb`) flag. We'll start with the following assembly language program (`debug.asm`) for this guide:

```
extern printf
global main

section .text
main:
    mov     rdi, format      ; argument #1
    mov     rsi, message     ; argument #2
    mov     rax, 0
    call    printf          ; call printf

    mov     rax, 0
    ret                     ; return 0

section .data
message:    db "Hello, world!", 0
format:     db "%s", 0xa, 0
courseCode: dq 2050
```

Below, we assemble and link the program (with the appropriate flags) so that we can use `gdb` on the resulting executable:

```
$ yasm -a x86 -m amd64 -g dwarf2 -f elf64 -o debug.o debug.asm
$ gcc -m64 -no-pie -o debug.out debug.o
$ gdb debug.out
```

We are now debugging our application.

## Part 2 – Layout Commands

In its most basic form, `gdb` isn't the most user-friendly. Layouts can make it easier by showing panels with useful information.

To keep the source code listing visible at all times, you can use the `layout asm` command. In order to display that source code in the syntax that we've been using, first set the assembly flavour to Intel:

```
(gdb) set disassembly-flavor intel
(gdb) layout asm
```

To keep a table of all of the registers visible at all times, you can use the `layout reg` command:

```
(gdb) layout reg
```

## Part 3 – Execution Commands

Like most debuggers, `gdb` will let us run our program, stopping at breakpoints, and even step through our program line-by-line. To set a breakpoint at the start of the `main` function, use the `break` (or `b` for short) command:

```
(gdb) b main
Breakpoint 1 at 0x401130: file debug.asm, line 6.
```

You can set a breakpoint at any label in an assembly language program. You can also set a breakpoint at any line in the original source file:

```
(gdb) b debug.asm:8
Breakpoint 2 at 0x40113e: file debug.asm, line 8.
```

To see a list of breakpoints, we can use the `info breakpoints` (or `info b` for short) command:

```
(gdb) info b
Num      Type           Disp Enb Address            What
1        breakpoint     keep y   0x0000000000401130 debug.asm:6
2        breakpoint     keep y   0x000000000040113e debug.asm:8
```

We can now run our program with the `run` (or `r` for short) command:

```
(gdb) r
Starting program:
/mnt/d/Winter2023/CSCI2050U/LectureExamples/assembly_language/using_gdb/debug.o
ut
```

```
Breakpoint 1, main () at debug.asm:6
```

Notice that `gdb` stopped at our breakpoint. We can continue running until the next breakpoint (line 8) using the `continue` (or `c` for short) command:

```
(gdb) c
Continuing.
```

```
Breakpoint 2, main () at debug.asm:8
```

We can also step through our program one instruction at a time, using the command `nexti` (or `ni` for short):

```
(gdb) ni
(gdb) ni
```

```
Hello, world!
```

The `layout asm` layout always shows us the source code, but we can see it without this layout using the `list` (or `l` for short) command if we don't want this layout:

```
(gdb) list
11      mov rax, 0
12      push rbx
13      call printf
14      pop rbx
15
16      mov rax, 0
17      call exit
18
19      section .data
20      format      db "%s", 0ah, 0dh, 0
```

## Part 4 – Data Commands

When debugging our program, we are probably going to want to know the state of our registers and variables. Without being able to do so, it will be challenging for us to identify where logic errors happen in our program. The easiest way to view the contents of a variable is using the `print` (or `p` for short) command:

```
(gdb) p (long)courseCode
$1 = 2050
```

We can also view register values this way:

```
(gdb) p (long)$rax
$2 = 14
```

It is also possible to print using type specifiers, and control how your output is displayed. A comprehensive set of type specifiers is given in the table, below:

Specifier	Meaning
t	binary (base [t]wo)
o	[o]ctal
x	he[x]adecimal
a	[a]ddress (hexadecimal absolute, plus hexadecimal offset from a close label)
c	[c]haracter
s	[s]tring
d	signed [d]ecimal
u	[u]nsigned decimal
f	[f]loating point

Examples of usage:

```
(gdb) p/x (long)courseCode
$3 = 0x802
```

```
(gdb) p/t (long)courseCode
$4 = 1000000000010
```

There is also the `x` (`e[x]amine`) command for viewing memory contents. This is useful for strings and arrays:

```
(gdb) x/s &message
0x404030:      "Hello, world!"
```

The `&` in the above command has the same meaning as in C/C++: “the address of”. This command has options similar to the `print` command. In general, the format of the command is:

x/nfu address

- n – how many of each data unit
- f – what type specifier (same as with `print`, but `i` is also possible for instructions)
- u – unit (data unit size)

Data unit sizes are given in the table below:

<b>Data Unit Size</b>	<b>Meaning</b>
b	[b]ytes
h	[h]alf words (words in x64 parlance)
w	[w]ords (double words or dwords in x64 parlance)
g	[g]iant words (quad words or qwords in x64 parlance)

Sample usage:

[illegible]

Another thing we are likely to want to do is to view the registers. This is not helpful if you have enabled register layout (`layout reg`), but is handy if you don't want to see the registers all the time. You can view the normal (integer) registers with the `info registers` command, and the floating point registers with the `info float` command:

```
(gdb) info r
rax          0xe          14
rbx          0x0          0
rcx          0xd          13
rdx          0x7ffff7dd59e0 140737351866848
rsi          0x7fffffff2      2147483634
rdi          0x1           1
rbp          0x0          0x0
rsp          0x7fffffffda88 0x7fffffffda88
```

```

r8          0xffffffff      4294967295
r9          0x0             0
r10         0x7ffff7dd26a0  140737351853728
r11         0x246           582
r12         0x400440        4195392
r13         0x7ffff7ffdb60  140737488345952
r14         0x0             0r15             0x0             0
rip         0x40054e        0x40054e <main+30>
eflags      0x212          [ AF IF ]
cs          0x33           51
ss          0x2b           43
ds          0x0            0
es          0x0            0
fs          0x0            0
gs          0x0            0

```

Finally, it might be useful to know how to quit gdb:

```

(gdb) quit
A debugging session is active.

    Inferior 1 [process 7900] will be killed.

Quit anyway? (y or n) y

```

## References

[1] <https://linux.die.net/man/1/gdb>